

SalientGS: Supplementary Material

Tianyu Xiong

School of Computer Science, Northwestern
Polytechnical University
Xi'an, China

Suning Ge

School of Computer Science, Northwestern
Polytechnical University
Xi'an, China

Rui Li

CEMSE, King Abdullah University of Science and
Technology
Thuwal, Saudi Arabia

Jiaqi Yang

School of Computer Science, Northwestern
Polytechnical University
Xi'an, China

1 Overview

This supplementary material accompanies the main paper “*SalientGS: Unified SfM-to-3DGS with Importance-Guided MCMC Gaussian Allocation.*” We provide:

- A released-code cross-benchmark verification with an explicit macro-average protocol (§2).
- A large-scene benchmark comparison (§3).
- Full SfM runtime tables for the Courthouse scene (§4).
- Full per-scene qualitative comparisons on all nine Mip-NeRF 360 scenes (§5).
- Per-image LPIPS analysis across ablation variants (§6).
- Threshold-sensitivity sweeps for τ_{imp} and the number of scoring views K (§7).
- Analysis of the SalientGS front-end: track-length distributions and a controlled track-cap ablation establishing that sufficient multi-view track coverage is a prerequisite for high reconstruction quality (§8).

2 Released-Code Cross-Benchmark Verification

Table 1 macro-averages the three dataset-level entries in the main paper, assigning equal weight to Mip-NeRF 360, Deep Blending, and Tanks & Temples. This is a cross-dataset summary, not a replacement for the individual benchmark results. Failed scenes remain in the corresponding dataset aggregate; notably, the *drjohnson* failures of GloSplat-A and VGGT-X are included rather than discarded. Under this explicit protocol, SalientGS has the best macro-average quality and runtime in the comparison. Its advantage reflects consistency across datasets: it is the only method with a top-two LPIPS result on every benchmark.

The released per-scene CSV uses seed 42, exactly 30K optimization steps, and a 1.5M Gaussian cap. All 13 runs execute feature extraction, matching, SfM, and training; reference sparse models are not used. For Mip-NeRF 360, the scene-equal aggregate used in the main comparison is 28.822 / 0.853 / 0.148 (PSNR / SSIM / LPIPS). Pooling all held-out images instead gives 29.607 / 0.874 / 0.140 because scenes contain different numbers of test images. We report the latter only as a diagnostic and never compare it against scene-averaged baselines.

Table 1: Cross-benchmark macro-average of the three dataset-level entries in the main comparison. Time is end-to-end minutes; each benchmark receives equal weight.

Method	Time↓	PSNR↑	SSIM↑	LPIPS↓
3DGS	28.35	26.98	0.855	0.211
3DGS-MCMC	28.85	27.40	0.872	0.191
Mini-Splatting	24.37	26.92	0.857	0.214
Speedy-splat	21.15	26.57	0.832	0.270
Taming-3DGS	14.71	26.96	0.840	0.251
DashGaussian	15.93	27.13	0.859	0.214
FastGS-big	13.54	27.48	0.861	0.211
GloSplat-A	22.01	23.15	0.750	0.265
VGGT-X	64.02	22.60	0.741	0.287
SalientGS	10.62	27.65	0.876	0.147

3 Large-Scene Benchmark Comparison

Table 3 reports SSIM, PSNR, and LPIPS on Building, Rubble, Residence, and Sci-Art; each competing method is cited in the first column. Our method achieves the best PSNR on Building, Rubble, and Residence, the best LPIPS on Building, Residence, and Sci-Art, and the best SSIM on Rubble, indicating that the unified SfM-to-3DGS pipeline remains highly competitive even on challenging large-scale scenes. Under the same large-scene training setup, SalientGS averages only about 1 h wall-clock training and ~5 GB peak GPU memory, compared to about 3 h and ~14 GB for CityGS-X [5].

4 SfM Runtime Details

Table 4 reports the full SfM component runtime on the Courthouse scene with 250–1000 input images, including GloSplat-F (XFeat+LightGlue with retrieval top-5; see §8 for details). These numbers correspond to the trend summarized by Figure 5 in the main paper.

5 Per-Scene Qualitative Comparisons

We present qualitative comparisons on all nine Mip-NeRF 360 [1] scenes. Each figure shows the rendered output from SalientGS (Ours) alongside ablation variants (Vanilla MCMC, Standard ADC, Freeze Poses) and a baseline method, with zoomed crops at the region of maximum LPIPS difference.

Table 2: Released-code verification by scene. Time includes feature extraction, matching, SfM, and 30K-step training on one RTX PRO 6000 Blackwell GPU. Every final model contains 1.5M Gaussians.

Dataset	Scene	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$	Time (min) $\downarrow$
Mip-NeRF 360	bicycle	27.125	0.828	0.135	10.37
	bonsai	33.207	0.955	0.106	13.53
	counter	29.604	0.919	0.136	14.03
	flowers	22.826	0.661	0.274	10.48
	garden	29.188	0.894	0.074	10.33
	kitchen	32.786	0.943	0.076	13.57
	room	31.945	0.928	0.150	12.72
	stump	28.310	0.836	0.123	10.86
Deep Blending	treehill	24.407	0.714	0.260	10.20
	drjohnson	28.824	0.892	0.210	10.69
Tanks & Temples	playroom	30.153	0.920	0.156	9.39
	train	23.003	0.846	0.128	10.32
	truck	26.306	0.893	0.090	9.73

Table 3: Large-scene benchmark comparison. Metrics are SSIM $\uparrow$, PSNR $\uparrow$, and LPIPS $\downarrow$. Red / orange / yellow denote the best / second-best / third-best values in each column.

Method	Building			Rubble			Residence			Sci-Art		
	SSIM	PSNR	LPIPS	SSIM	PSNR	LPIPS	SSIM	PSNR	LPIPS	SSIM	PSNR	LPIPS
<i>w/ Geometric Optimization</i>												
SuGaR [6]	0.507	17.76	0.455	0.577	20.69	0.453	0.603	18.74	0.406	0.698	18.60	0.349
NeuS [15]	0.463	18.01	0.611	0.480	20.46	0.618	0.503	17.85	0.533	0.633	18.62	0.472
Neuralangelo [9]	0.582	17.89	0.322	0.625	20.18	0.314	0.644	18.03	0.263	0.769	19.10	0.231
PGSR [2]	0.480	16.12	0.573	0.728	23.09	0.334	0.746	20.57	0.289	0.799	19.72	0.275
VCR-Gaus [3]	0.502	19.56	0.502	0.541	21.34	0.428	0.623	20.59	0.359	0.665	19.31	0.465
CityGaussianV2 [12]	0.650	19.07	0.397	0.720	23.75	0.322	0.769	21.15	0.234	0.810	20.66	0.266
UrbanGS [8]	0.802	22.82	0.208	0.791	26.25	0.210	0.823	22.48	0.205	0.824	22.62	0.279
CityGS-X [5]	0.817	22.76	0.191	0.823	26.15	0.210	0.819	22.44	0.194	0.867	22.77	0.179
<i>w/o Geometric Optimization</i>												
Mega-NeRF [14]	0.547	20.92	0.454	0.553	24.06	0.508	0.628	22.08	0.401	0.770	25.60	0.312
Switch-NeRF [13]	0.579	21.54	0.397	0.562	24.31	0.478	0.654	22.57	0.352	0.795	26.51	0.271
VastGaussian [10]	0.728	21.80	0.225	0.742	25.20	0.264	0.699	21.01	0.261	0.761	22.64	0.261
3DGS [7]	0.738	22.53	0.214	0.725	25.51	0.316	0.745	22.36	0.247	0.791	24.13	0.262
DoGaussian [4]	0.759	22.73	0.204	0.765	25.78	0.257	0.740	21.94	0.244	0.804	24.42	0.219
CityGaussian [11]	0.778	21.55	0.246	0.813	25.77	0.228	0.813	22.00	0.211	0.837	21.39	0.230
SalientGS (Ours)	0.804	24.19	0.181	0.824	26.98	0.213	0.816	22.67	0.193	0.857	22.91	0.173

Table 4: SfM runtime comparison (seconds). We compare the SfM component runtime across different methods on the Courthouse scene with varying numbers of input images. SalientGS combines Fisher Vector retrieval with first-order SfM optimization to achieve significant speedups over traditional and learning-based approaches.

# Images	COLMAP	MASt3R	GloSplat-A	VGGT-X	GloSplat-F	SalientGS
250	644.6	312.2	222.2	53.9	100.2	47.8
500	1165.0	638.0	408.2	113.3	191.9	98.0
750	2226.3	978.5	855.7	249.6	250.0	138.5
1000	4349.2	1308.2	1371.3	398.9	327.8	186.1

Outdoor Scenes. Figures 1–5 cover the five outdoor scenes. The crops show localized differences in foliage, bark,

and grass, while Table 5 provides the scene-level ranking. SalientGS is best on *bicycle* and *stump* and ties Vanilla MCMC on *treehill* after rounding. On *flowers*, it trails Vanilla MCMC (0.274 vs. 0.264) but remains better than Standard ADC and Freeze Poses; on *garden*, Standard ADC is best (0.066 vs. 0.074 for SalientGS). Thus the qualitative crops illustrate local behavior rather than overriding the per-scene metrics.

Indoor Scenes. Figures 6–9 cover the four indoor scenes. SalientGS leads all displayed ablation variants on *bonsai*, *kitchen*, and *room*. On *counter*, it records 0.136 LPIPS and therefore trails Vanilla MCMC (0.132) and Freeze Poses (0.134), while remaining better than Standard ADC (0.148). The selected *room* failure view is a separate per-view diagnostic: SalientGS trails VGGT-X there (0.312 vs. 0.244),

Table 5: Mean per-image LPIPS ($\downarrow$) on Mip-NeRF 360 scenes. Scene rows show per-scene means. We report both the scene-equal mean used in the main comparison and the pooled test-image diagnostic. Best displayed result per scene in bold.

Scene	SalientGS	Vanilla MCMC	Std. ADC	Freeze Poses
<i>Outdoor Scenes</i>				
bicycle	0.135	0.152	0.153	0.173
flowers	0.274	0.264	0.309	0.293
garden	0.074	0.074	0.066	0.081
stump	0.123	0.140	0.145	0.151
treehill	0.260	0.260	0.269	0.270
<i>Indoor Scenes</i>				
bonsai	0.106	0.169	0.171	0.170
counter	0.136	0.132	0.148	0.134
kitchen	0.076	0.236	0.167	0.173
room	0.150	0.271	0.252	0.271
Scene mean	0.148	0.189	0.187	0.192
Pooled images	0.140	0.189	0.187	0.191

despite its 0.150 scene-mean LPIPS and its lead over the displayed ablation variants.

6 Per-Image LPIPS Analysis

Table 5 reports the corrected full-model evaluation together with the verified ablation measurements under the same metric protocol. The earlier discrepancy was an evaluation error in the full-model aggregate, not a change of experimental protocol. At full precision, SalientGS improves over vanilla MCMC on seven scenes and trails it on *flowers* and *counter*; *garden* and *treehill* appear tied after rounding. Its nine-scene mean is substantially lower (0.148 vs. 0.189). Scene-level exceptions are retained explicitly rather than hidden by the aggregate.

7 Threshold Sensitivity Sweep

Table 6 reports the Mip-NeRF 360 sensitivity sweep referenced in the main paper. We vary the importance threshold τ_{imp} and the number of views K used for multi-view score computation while keeping the rest of the pipeline fixed.

The released setting uses $\tau_{\text{imp}}=5$ and $K=10$. These defaults were fixed before the full benchmark and are not selected per scene from the sweep. PSNR and SSIM vary mildly across nearby settings, while LPIPS is more sensitive to perceptual blur and high-frequency detail; we therefore report the three metrics separately rather than summarizing robustness with a single PSNR range.

8 Front-End Track Analysis and Track-Length Ablation

A natural question is whether the quality advantage of SalientGS over retrieval-only front-ends (such as GloSplat-F, which uses XFeat+LightGlue with top-5 retrieval) originates primarily from better Gaussian allocation or from a richer multi-view track structure. We investigate this through two complementary experiments: (1) a comparison of the 3D-point track-length distributions produced by

Table 6: Threshold-sensitivity sweep on Mip-NeRF 360. We vary the importance threshold τ_{imp} and the number of scoring views K . The selected full-model configuration is marked as *default*.

Group	Value	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$
τ_{imp}	1	28.688	0.8444	0.159
	3	28.685	0.8448	0.157
	5 (<i>default</i>)	28.822	0.8531	0.148
	10	28.732	0.8466	0.153
	20	28.759	0.8477	0.151
K	3	28.696	0.8457	0.156
	5	28.689	0.8455	0.156
	10 (<i>default</i>)	28.822	0.8531	0.148
	15	28.714	0.8450	0.157
	20	28.679	0.8446	0.157

three front-ends on the *bicycle* scene, and (2) a controlled track-length cap ablation that holds the front-end fixed and varies only the maximum number of observations retained per track before global positioning and bundle adjustment.

Track-Length Distributions. Table 7 compares the track-length distributions produced on the *bicycle* scene by three pipelines: GloSplat-A (SIFT exhaustive), GloSplat-F (XFeat+LightGlue, retrieval top-5), and SalientGS (SIFT + Fisher Vector retrieval, $k=20$). SalientGS reconstructs $\sim 4\%$ fewer 3D points than GloSplat-A (51.8k vs. 54.3k), but its track-length distribution remains very close: median and max are nearly identical, the mean is only marginally lower (4.47 vs. 4.69), and all four length-bucket percentages differ by at most 0.6 pp. In contrast, GloSplat-F—which uses only 5 retrieval neighbors—yields $\sim 16\%$ fewer points with a substantially shorter track distribution (median 3 vs. 4, max 20 vs. 74), and long tracks (≥ 9 observations) are nearly absent (2.1% vs. 9.5%). The MST connectivity guarantee built into SalientGS’s Fisher Vector retrieval (Sec. 3.1 of the main paper) is responsible for preserving the exhaustive-like track structure: the $k=20$ candidate graph, augmented with MST edges, achieves far denser coverage than a bare top-5 list.

Table 7: Track-length distribution on the bicycle scene for GloSplat-A (SIFT exhaustive), GloSplat-F (XFeat+LightGlue, retrieval top-5), and SalientGS (Fisher Vector, $k=20$).


	GloSplat-A	GloSplat-F	SalientGS
Total points	54,275	45,507	51,847
Mean / Median	4.69 / 4	3.75 / 3	4.47 / 4
Max	74	20	71
2–3	26,861 (49.5%)	27,182 (59.7%)	25,814 (49.8%)
4–8	22,036 (40.6%)	17,366 (38.2%)	21,124 (40.7%)
9–16	4,694 (8.6%)	948 (2.1%)	4,133 (8.0%)
17+	684 (1.3%)	11 (0.0%)	776 (1.5%)

Figure 1: Per-scene qualitative comparison: *bicycle*.Figure 2: Per-scene qualitative comparison: *flowers*.

Track-Length Cap Ablation. To establish causality between track length and downstream rendering quality, we conduct a controlled ablation: we fix each method’s front-end (i.e., the same pair graph and matches) and vary only the maximum number of observations retained per track before global positioning and bundle adjustment. Table 8 reports the effect on Mip-NeRF 360 (averaged over 9 scenes). The $L=2$ condition reduces every track to a pairwise observation—simulating a worst-case version of the track shortening seen in GloSplat-F—while $L=all$ retains the full distribution. For GloSplat (SIFT exhaustive front-end), capping at $L=2$ causes a 1.40 dB PSNR drop ($28.86 \rightarrow 27.46$),

confirming that long multi-view tracks provide substantial geometric constraints that improve triangulation robustness and stabilize bundle adjustment. In the matched track-cap experiment, SalientGS exhibits a similar degradation of 1.29 dB ($28.82 \rightarrow 27.53$) under the same $L=2$ cap. Once every track is reduced to a pairwise observation, both pipelines operate on geometrically equivalent constraints regardless of their front-end or allocation strategy, leaving the underlying bundle adjustment quality as the dominant factor.

Taken together, these results support two conclusions. First, SalientGS’s Fisher Vector retrieval with $k=20$ and MST connectivity is sufficient to recover an exhaustive-quality

Figure 3: Per-scene qualitative comparison: *garden*.Figure 4: Per-scene qualitative comparison: *stump*.

track distribution at a fraction of the matching cost, closing the structural gap that limits GloSplat-F. Second, the near-identical $L=2$ scores across both methods confirm that the gain from full tracks is a property of multi-view triangulation geometry, not of any particular front-end or allocation design—making rich track coverage a necessary condition for high-quality reconstruction.


9 Additional Camera-Ready Ablations

This section records the ablations completed during the author response. A previous evaluation error affected only the reported full-model aggregate. All occurrences of the

complete default model therefore use the corrected value, 28.82 dB / 0.853 / 0.148; the other variant measurements were verified under the same protocol and remain unchanged.

Table 9 shows that guided allocation improves both initializers: LPIPS changes from 0.151 to 0.133 for GLOMAP and from 0.149 to 0.148 for FastMap. This supports an allocator contribution that is not specific to one camera initializer.

The sensitivity tables expose the cost-quality trade-off rather than retroactively selecting the best row. For example, $T=200$ improves the three reported metrics but computes scores $2.5\times$ as often as the released $T=500$ schedule;

Figure 5: Per-scene qualitative comparison: *treehill*.Figure 6: Per-scene qualitative comparison: *bonsai*.

alternative selection quantiles and opacity mixing can also improve LPIPS. We retain the preselected defaults in the full benchmark to avoid per-scene or post-hoc configuration selection.

10 Additional Robustness and Failure Diagnostics

The diagnostics in Table 17 identify the texture-poor *room* case through collapsed track length, greater MST reliance, and a high early residual. A feed-forward fallback improves the failure-view LPIPS from 0.312 to 0.235, slightly better

than the VGGT-X reference of 0.244. Sparse/few-view reconstruction and severe repetitive or symmetric appearance remain limitations; learned retrieval and adaptive geometry-aware budgets are natural future extensions.



Figure 7: Per-scene qualitative comparison: *counter*.



Figure 8: Per-scene qualitative comparison: *kitchen*.



Figure 9: Qualitative comparison on *room*. The selected failure view exposes a texture-poor local region where SalientGS trails VGGT-X; this per-view observation is distinct from the 0.150 scene-mean LPIPS in Table 5.

Table 8: Track-length cap ablation on MipNeRF360 (9-scene avg.). The front-end is fixed for each method; only the maximum retained observations per track varies.

Max track length	PSNR $\uparrow$	SSIM $\uparrow$	LPIPS $\downarrow$
<i>GloSplat (SIFT exhaustive front-end)</i>			
<i>GloSplat-F (retrieval top-5, ref.)</i>	27.77	0.818	0.164
$L = 2$	27.46	0.804	0.174
$L = \text{all}$	28.86	0.862	0.139
<i>GloSplat-A (exhaustive, ref.)</i>	28.86	0.862	0.139
<i>SalientGS (Fisher Vector $k=20$ front-end)</i>			
$L = 2$	27.53	0.808	0.172
$L = \text{all (full model)}$	28.82	0.853	0.148

Table 9: SfM initialization versus Gaussian allocator on Mip-NeRF 360. All rows use the same joint training protocol.

SfM initialization	Allocator	PSNR↑	SSIM↑	LPIPS↓
GLOMAP	ADC	27.52	0.827	0.168
GLOMAP	Vanilla MCMC	28.70	0.846	0.151
GLOMAP	Guided MCMC	28.87	0.855	0.133
FastMap	ADC	27.67	0.828	0.162
FastMap	Vanilla MCMC	28.72	0.848	0.149
FastMap	Guided MCMC (full model)	28.82	0.853	0.148

Table 10: Importance-score schedule and opacity mixing on Mip-NeRF 360. Bold marks the best value within each sub-sweep; default/full identifies the fixed released configuration, not a post-hoc metric optimum.

Setting	PSNR↑	SSIM↑	LPIPS↓
$T = 100$	28.83	0.855	0.135
$T = 200$	28.87	0.856	0.133
$T = 500$ (default/full)	28.82	0.853	0.148
$T = 1000$	28.80	0.852	0.137
score at iter. 0	28.78	0.851	0.138
score at iter. 3000 (default/full)	28.82	0.853	0.148
$\lambda_{\text{mix}} = 0$	28.78	0.851	0.141
$\lambda_{\text{mix}} = 0.05$ (default/full)	28.82	0.853	0.148
$\lambda_{\text{mix}} = 0.20$	28.75	0.850	0.145

Table 11: Robust quantile normalization and selection-quantile interaction on Mip-NeRF 360. Bold marks the best value by metric; default/full identifies the fixed released configuration. The final row reports the fraction of Gaussians in the high-error mass at 30K iterations.

Setting	PSNR↑	SSIM↑	LPIPS↓
Plain min-max	28.71	0.850	0.150
Default quantiles (full)	28.82	0.853	0.148
$(q_{\text{hi}}, q_{\text{lo}}) = (.80, .20)$	28.79	0.852	0.139
$(q_{\text{hi}}, q_{\text{lo}}) = (.95, .05)$	28.82	0.851	0.137
High-error mass at 30K	robust: 9.1%; min-max: 42.5%		

Table 12: Pose error before and after joint optimization on the ETH3D SLAM evaluation. Errors are in centimeters; percentages are fractions of cameras below the indicated threshold.

Method	Median↓	Mean↓	< 5cm↑	< 10cm↑
FastMap init	9.8	18.5	31.6	50.5
FastMap + joint	6.9	13.5	38.2	57.0
GLOMAP init	8.5	14.6	35.5	55.5
GLOMAP + joint	7.1	12.6	37.8	56.5

Table 13: Per-scene Deep Blending results. The competing methods fail on *drjohnson*, which explains the averaged values in the main table.

Scene	Method	PSNR↑	SSIM↑	LPIPS↓
<i>playroom</i>	GloSplat-A	29.38	0.903	0.224
<i>playroom</i>	SalientGS	30.15	0.920	0.156
<i>drjohnson</i>	GloSplat-A	7.52	0.263	0.792
<i>drjohnson</i>	VGGT-X	7.65	0.337	0.858
<i>drjohnson</i>	SalientGS	28.82	0.892	0.210

Table 14: Budget efficiency on Mip-NeRF 360.

N_{GS}	Vanilla MCMC	Guided MCMC	ΔPSNR
500K	28.19	28.46	+0.27
1.0M	28.59	28.81	+0.22
1.5M	28.72	28.82	+0.10
2.0M	28.80	28.93	+0.13
3.0M	28.88	28.97	+0.09

Table 15: Dense indoor robustness. SG denotes SalientGS and V+M denotes the VGGT-X+MCMC comparison.

Scene	SG PSNR	V+M PSNR	SG LPIPS	V+M LPIPS
bonsai	33.21	30.71	0.106	0.139
counter	29.60	28.80	0.136	0.158
kitchen	32.79	29.68	0.076	0.113
room	31.95	30.76	0.150	0.165
Average	31.89	29.99	0.117	0.144

Table 16: Additional large-scene and drone-capture evidence.

Setting	Result
Mill-19 Building	Best PSNR, 24.19 dB
Mill-19 Rubble	Best PSNR, 26.98 dB
Four large scenes	Best PSNR/LPIPS on 3 of 4 scenes
Runtime/memory	~1 h / ~5 GB vs. ~3 h / ~14 GB for CityGS-X

Table 17: Failure diagnostics and fallback on representative Mip-NeRF 360 scenes. Residual is the normalized early epipolar residual.

Scene	Median track	Mean degree	MST reliance	Residual
garden	4	18.4	4%	0.20
kitchen	3	13.4	11%	0.43
room	2	9.7	21%	0.57

room LPIPS: original 0.312; fallback 0.235; VGGT-X 0.244

Table 18: Sequential Tanks & Temples comparison with CF-3DGS.

Scene	Method	PSNR↑	SSIM↑	LPIPS↓
Family	CF-3DGS	33.12	0.960	0.051
Family	SalientGS	36.80	0.970	0.024
Francis	CF-3DGS	32.78	0.915	0.135
Francis	SalientGS	36.29	0.947	0.056

References

- [1] Jonathan T Barron, Ben Mildenhall, Dor Verbin, Pratul P Srinivasan, and Peter Hedman. 2022. Mip-nerf 360: Unbounded anti-aliased neural radiance fields. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. IEEE, Piscataway, NJ, USA, 5470–5479.
- [2] Danpeng Chen, Hai Li, Weicai Ye, Yifan Wang, Weijian Xie, Shangjin Zhai, Nan Wang, Haomin Liu, Hujun Bao, and Guofeng Zhang. 2025. PGSR: Planar-Based Gaussian Splatting for Efficient and High-Fidelity Surface Reconstruction. *IEEE Transactions on Visualization and Computer Graphics* 31, 12 (2025), 6100–6111. doi:10.1109/TVCG.2024.3494046
- [3] Hanlin Chen, Fangyin Wei, Chen Li, Tianxin Huang, Yunsong Wang, and Gim Hee Lee. 2024. VCR-GauS: View Consistent Depth-Normal Regularizer for Gaussian Surface Reconstruction. In *Advances in Neural Information Processing Systems*, Vol. 37. Curran Associates, Inc., Red Hook, NY, USA.
- [4] Yu Chen and Gim Hee Lee. 2024. DOGS: Distributed-Oriented Gaussian Splatting for Large-Scale 3D Reconstruction Via Gaussian Consensus. In *Advances in Neural Information Processing Systems*, Vol. 37. Curran Associates, Inc., Red Hook, NY, USA.
- [5] Yuanyuan Gao, Hao Li, Jiaqi Chen, Zhengyu Zou, Zhihang Zhong, Dingwen Zhang, Xiao Sun, and Junwei Han. 2025. CityGS-X: A Scalable Architecture for Efficient and Geometrically Accurate Large-Scale Scene Reconstruction. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*. IEEE, Piscataway, NJ, USA.
- [6] Antoine Guédon and Vincent Lepetit. 2024. SuGaR: Surface-Aligned Gaussian Splatting for Efficient 3D Mesh Reconstruction and High-Quality Mesh Rendering. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. IEEE, Piscataway, NJ, USA, 5354–5363.
- [7] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkühler, and George Drettakis. 2023. 3D Gaussian Splatting for Real-Time Radiance Field Rendering. *ACM Transactions on Graphics* 42, 4, Article 139 (July 2023), 14 pages. <https://repo-sam.inria.fr/fungraph/3d-gaussian-splatting/>
- [8] Changbai Li, Haodong Zhu, Hanlin Chen, Xiuping Liang, Tongfei Chen, Shuwei Shao, Linlin Yang, Huobin Tan, and Baochang Zhang. 2026. UrbanGS: Efficient and Scalable Architecture for Geometrically Accurate Large-Scale Scene Reconstruction. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=L3utaw6SD9>
- [9] Zhaoshuo Li, Thomas Müller, Alex Evans, Russell H. Taylor, Mathias Unberath, Ming-Yu Liu, and Chen-Hsuan Lin. 2023. Neuralangelo: High-Fidelity Neural Surface Reconstruction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. IEEE, Piscataway, NJ, USA, 8456–8465.
- [10] Jiaqi Lin, Zhihao Li, Xiao Tang, Jianzhuang Liu, Shiyong Liu, Jiayue Liu, Yangdi Lu, Xiaofei Wu, Songcen Xu, Youliang Yan, and Wenming Yang. 2024. VastGaussian: Vast 3D Gaussians for Large Scene Reconstruction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. IEEE, Piscataway, NJ, USA.
- [11] Yang Liu, He Guan, Chuanchen Luo, Lue Fan, Naiyan Wang, Junran Peng, and Zhaoxiang Zhang. 2024. CityGaussian: Real-time High-quality Large-Scale Scene Rendering with Gaussians. In *European Conference on Computer Vision*. Springer, Cham.
- [12] Yang Liu, Chuanchen Luo, Zhongkai Mao, Junran Peng, and Zhaoxiang Zhang. 2025. CityGaussianV2: Efficient and Geometrically Accurate Reconstruction for Large-Scale Scenes. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=a3ptUbuzbW>
- [13] Zhenxing Mi and Dan Xu. 2023. Switch-NeRF: Learning Scene Decomposition with Mixture of Experts for Large-scale Neural Radiance Fields. In *International Conference on Learning Representations*. <https://openreview.net/forum?id=PQ2zoIZqvm>
- [14] Haithem Turki, Deva Ramanan, and Mahadev Satyanarayanan. 2022. Mega-NeRF: Scalable Construction of Large-Scale NeRFs for Virtual Fly-Throughs. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. IEEE, Piscataway, NJ, USA, 12922–12931.
- [15] Peng Wang, Lingjie Liu, Yuan Liu, Christian Theobalt, Taku Komura, and Wenping Wang. 2021. NeuS: Learning Neural Implicit Surfaces by Volume Rendering for Multi-view Reconstruction. In *Advances in Neural Information Processing Systems*, Vol. 34.

Curran Associates, Inc., Red Hook, NY, USA.